

Bayesian Learning

Lecture 9 - HMC and Stan

Bertil Wegmann

Department of Computer and Information Science
Linköping University



Lecture overview

- Hamiltonian Monte Carlo
- Stan

Hamiltonian Monte Carlo

- When $\theta = (\theta_1, \dots, \theta_p)$ is **high-dimensional**, $p(\theta|y)$ usually located in some subregion of \mathbb{R}^p with complicated geometry.
- MH: hard to find good proposal distribution $q(\cdot|\theta^{(i-1)})$.
- MH: use very small step sizes otherwise too many rejections.
- **Hamiltonian Monte Carlo (HMC)**:
 - ▶ distant proposals **and**
 - ▶ high acceptance probabilities.
- HMC: add extra **momentum** parameters $\phi = (\phi_1, \dots, \phi_p)$ and sample from

$$p(\theta, \phi|y) = p(\theta|y) p(\phi)$$

Hamiltonian Monte Carlo

- Physics: **Hamiltonian** system $H(\theta, \phi) = U(\theta) + K(\phi)$, where U is the **potential energy** and K is the **kinetic energy**.
- **Hamiltonian Dynamics**

$$\begin{aligned}\frac{d\theta_i}{dt} &= \frac{\partial H}{\partial \phi_i} = \frac{\partial K}{\partial \phi_i}, \\ \frac{d\phi_i}{dt} &= -\frac{\partial H}{\partial \theta_i} = -\frac{\partial U}{\partial \theta_i}\end{aligned}$$

- Hockey puck sliding over a friction-less surface: [illustration](#).
- Use $U(\theta) = -\log[p(\theta)p(y|\theta)]$.
- Use $\phi \sim N(0, M)$ where M is the mass matrix and

$$K(\phi) = -\log[p(\phi)] = \frac{1}{2}\phi^T M^{-1}\phi + \text{const}$$

Hamiltonian Monte Carlo

■ Hamiltonian Dynamics

$$\begin{aligned}\frac{d\theta_i}{dt} &= [M^{-1}\phi]_i, \\ \frac{d\phi_i}{dt} &= \frac{\partial \log p(\theta|y)}{\partial \theta_i}\end{aligned}$$

which can be simulated using the **leapfrog algorithm**

$$\begin{aligned}\phi_i\left(t + \frac{\varepsilon}{2}\right) &= \phi_i(t) + \frac{\varepsilon}{2} \frac{\partial \log p(\theta(t)|y)}{\partial \theta_i}, \\ \theta(t + \varepsilon) &= \theta(t) + \varepsilon M^{-1}\phi(t), \\ \phi_i(t + \varepsilon) &= \phi_i\left(t + \frac{\varepsilon}{2}\right) + \frac{\varepsilon}{2} \frac{\partial \log p(\theta(t)|y)}{\partial \theta_i},\end{aligned}$$

where ε is the step size.

■ **Discretization** \Rightarrow acceptance probability drops with ε .

The Hamiltonian Monte Carlo algorithm

■ Initialize $\theta^{(0)}$ and iterate for $i = 1, 2, \dots$

- 1 Sample the starting **momentum** $\phi_s \sim N(0, M)$
- 2 Simulate new values for (θ_p, ϕ_p) by iterating the **leapfrog algorithm** L times, starting in $(\theta^{(i-1)}, \phi_s)$.
- 3 Compute the **acceptance probability**

$$\alpha = \min \left(1, \frac{p(y|\theta_p)p(\theta_p)}{p(y|\theta^{(i-1)})p(\theta^{(i-1)})} \frac{p(\phi_p)}{p(\phi_s)} \right)$$

- 4 With probability α set $\theta^{(i)} = \theta_p$ and $\phi^{(i)} = \phi_p$ otherwise.

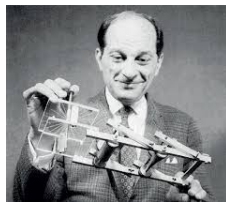
■ **Tuning parameters:** 1. stepsize ε , 2. number of leapfrog iterations L and 3. mass matrix M . **No U-turn.**

Stan

- **Stan** is a probabilistic programming language based on HMC.
- Allows for Bayesian inference in many models with automatic implementation of the MCMC sampler.
- Named after Stanislaw Ulam (1909-1984), co-inventor of the Monte Carlo algorithm.
- Written in C++ but can be run from R using the package `rstan`




Stan logo



Stanislaw Ulam

Stan - toy example: three plants

- Three plants were observed for four months, measuring the number of flowers



Stan Model 1: iid normal

$$y_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$$

```
library(rstan)
y=c(4,5,6,4,0,2,5,3,8,6,10,8)
N=length(y)

StanModel = '
data {
  int<lower=0> N; // Number of observations
  int<lower=0> y[N]; // Number of flowers
}
parameters {
  real mu;
  real<lower=0> sigma2;
}
model {
  mu ~ normal(0,100); // Normal with mean 0, st.dev. 100
  sigma2 ~ scaled_inv_chi_square(1,2); // Scaled-inv-chi2 with nu 1, sigma 2
  for(i in 1:N){
    y[i] ~ normal(mu,sqrt(sigma2));
  }
},'
```

Stan Model 2: multilevel normal

$$y_{t,p} \sim N\left(\mu_p, \sigma_p^2\right), \mu_p \sim N\left(\mu, \sigma^2\right)$$

```
StanModel <- '  
data {  
  int<lower=0> N; // Number of observations  
  int<lower=0> y[N]; // Number of flowers  
  int<lower=0> P; // Number of plants  
}  
transformed data {  
  int<lower=0> M; // Number of months  
  M = N / P;  
}  
parameters {  
  real mu;  
  real<lower=0> sigma2;  
  real mup[P];  
  real sigmap2[P];  
}  
model {  
  mu ~ normal(0,100); // Normal with mean 0, st.dev. 100  
  sigma2 ~ scaled_inv_chi_square(1,2); // Scaled-inv-chi2 with nu 1, sigma 2  
  for(p in 1:P){  
    mup[p] ~ normal(mu,sqrt(sigma2));  
    for(m in 1:M) {  
      y[M*(p-1)+m] ~ normal(mup[p],sqrt(sigmap2[p]));  
    }  
  }  
}
```

Stan Model 3: multilevel Poisson

$$y_{t,p} \sim \text{Poisson}(\mu_p), \mu_p \sim \text{logN}(\mu, \sigma^2)$$

```
StanModel <- '
data {
  int<lower=0> N; // Number of observations
  int<lower=0> y[N]; // Number of flowers
  int<lower=0> P; // Number of plants
}
transformed data {
  int<lower=0> M; // Number of months
  M = N / P;
}
parameters {
  real mu;
  real<lower=0> sigma2;
  real mup[P];
}
model {
  mu ~ normal(0,100); // Normal with mean 0, st.dev. 100
  sigma2 ~ scaled_inv_chi_square(1,2); // Scaled-inv-chi2 with nu 1, sigma 2
  for(p in 1:P){
    mup[p] ~ lognormal(mu,sqrt(sigma2)); // Log-normal
    for(m in 1:M) {
      y[M*(p-1)+m] ~ poisson(mup[p]); // Poisson
    }
  }
}
```

Stan: fit model and analyze output

```
data <- list(N=N, y=y, P=P)
warmup <- 1000
niter <- 2000
fit <- stan(file="StanModel.stan", data=data, warmup=warmup, iter=niter, chains=4, cores=2)

# Print the fitted model
print(fit,digits_summary=3)

# Extract posterior samples
postDraws <- extract(fit)

# Do traceplots of the first chain
par(mfrow = c(1,1))
plot(postDraws$mu[1:(niter-warmup)],type="l",ylab="mu",main="Traceplot")

# Do automatic traceplots of all chains
traceplot(fit)

# Bivariate posterior plots
pairs(fit)
```

Stan - useful links

- [Getting started with RStan](#)
- [RStan vignette](#)
- [Stan Modeling Language User's Guide and Reference Manual](#)
- [Stan Case Studies](#)